

## 1. **salloc** command line used

To allocate a compute node interactively:

```
salloc --cpus-per-task=8 <executable_path>
```

## 2. Maximum number of cores M on the node

80 (I used the command )

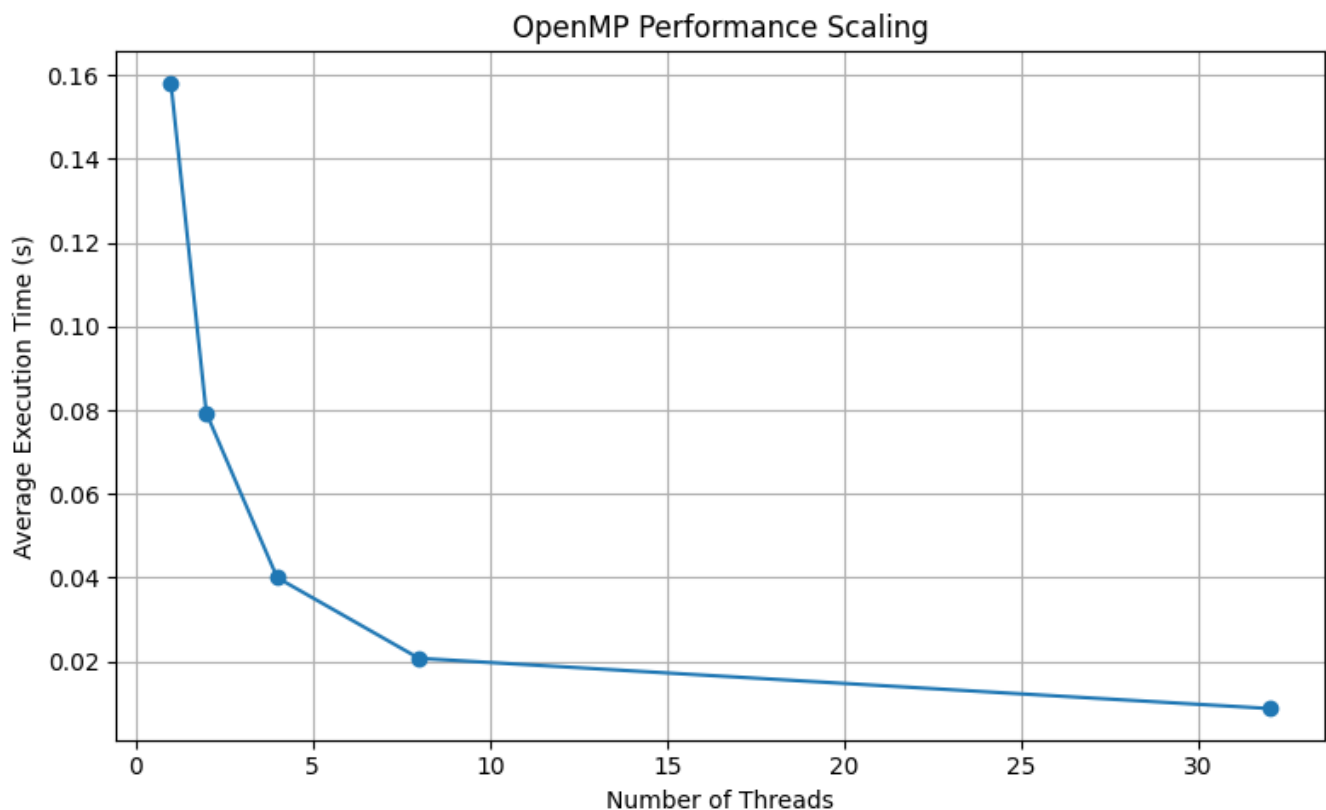
## 3. How I set different $N_t$ 's

```
export OMP_NUM_THREADS=1    # For N1
./omp_hello2

export OMP_NUM_THREADS=2    # For N2
./omp_hello2


export OMP_NUM_THREADS=4    # For N3
./omp_hello2
```

## 4. Record execution times and plot averages



## 5. Speedup relative to $N_1$

(each point represents  $N_i$  at  $i$ )



## 6. Questions

### (a) Why printed messages may differ in order for a given $N_t$

OpenMP threads run concurrently; the exact timing of each thread is not known in advance, printing from multiple threads can cause random-looking output orders.

### (b) Why execution times may be larger for larger $N_t$

Thread management overhead increases with more threads.

### (c) How to modify the code to obtain speedup $> 1$ for all $N_t$

By choosing values of  $N_t$  which lie between 16 and 64?

### (d) Type of parallelism in `omp_hello2.c`

Data parallelism

### (e) Modify code to avoid `if (tid == 0)` while keeping master print

```
#pragma omp parallel
{
    int tid = omp_get_thread_num();
    #pragma omp master
    {
        printf("Hello from the master thread!\n");
    }
}
```

```
    printf("Hello from thread %d\n", tid);  
}
```